

Keras -- MLPs on MNIST

```
In [1]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

```
In [0]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [0]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [4]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)

Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of
1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.sh
ape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2
])
```

```
In [6]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each imag
e is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image
is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
In [7]: # An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0
```

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
5	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	25	
4	247	127	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	15
0	170	253	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0	
2	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251	93	8
3	82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	18	219	25
0	253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	15
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	24
0	225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	18
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

```

0
0 0 0 0 0 0 0 249 253 249 64 0 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 46 130 183 25
3
253 207 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 0 0 0 39 148 229 253 253 253 250 182 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 0 0 24 114 221 253 253 25
3
253 201 78 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 0 23 66 213 253 253 253 253 198 81 2 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 18 171 219 253 253 253 253 19
5
80 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
55 172 226 253 253 253 253 244 133 11 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 136 253 253 253 212 135 132 1
6
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0
0 0 0 0 0 0 0 0 0 0 0]

```

```

In [0]: # if we observe the above matrix each cell is having a value between 0-255
        # before we move to apply machine learning algorithms lets try to normalize the data

```

```
X_train = X_train/255
X_test = X_test/255
```

[illegible]

0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686

[illegible]

```
In [10]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```


Softmax classifier

```
In [0]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances
# to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias=True)
```

```

as is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.
X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

```

In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

1.USING TWO HIDDEN LAYERS

a.MLP + ReLU + ADAM

```
In [62]: model_relu = Sequential()
model_relu.add(Dense(396, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.071, seed=None)))
model_relu.add(Dense(198, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=None) ))
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_44 (Dense)	(None, 396)	310860
dense_45 (Dense)	(None, 198)	78606
dense_46 (Dense)	(None, 10)	1990

=====
Total params: 391,456
Trainable params: 391,456
Non-trainable params: 0
=====
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 3s 49us/step - loss: 0.2327 - acc: 0.9309 - val_loss: 0.1019 - val_acc: 0.9675
Epoch 2/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0

```
865 - acc: 0.9739 - val_loss: 0.0915 - val_acc: 0.9716
Epoch 3/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
574 - acc: 0.9826 - val_loss: 0.0709 - val_acc: 0.9771
Epoch 4/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
386 - acc: 0.9879 - val_loss: 0.0716 - val_acc: 0.9777
Epoch 5/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
262 - acc: 0.9921 - val_loss: 0.0691 - val_acc: 0.9794
Epoch 6/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
207 - acc: 0.9935 - val_loss: 0.0786 - val_acc: 0.9773
Epoch 7/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
182 - acc: 0.9943 - val_loss: 0.0695 - val_acc: 0.9799
Epoch 8/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0
144 - acc: 0.9952 - val_loss: 0.0734 - val_acc: 0.9797
Epoch 9/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0
130 - acc: 0.9958 - val_loss: 0.0860 - val_acc: 0.9774
Epoch 10/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0
121 - acc: 0.9960 - val_loss: 0.0870 - val_acc: 0.9782
Epoch 11/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0
098 - acc: 0.9969 - val_loss: 0.0862 - val_acc: 0.9800
Epoch 12/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
081 - acc: 0.9974 - val_loss: 0.0941 - val_acc: 0.9780
Epoch 13/20
60000/60000 [=====] - 2s 32us/step - loss: 0.0
111 - acc: 0.9964 - val_loss: 0.0881 - val_acc: 0.9794
Epoch 14/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
108 - acc: 0.9963 - val_loss: 0.1193 - val_acc: 0.9754
Epoch 15/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
```

```

074 - acc: 0.9977 - val_loss: 0.0971 - val_acc: 0.9798
Epoch 16/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
089 - acc: 0.9968 - val_loss: 0.1021 - val_acc: 0.9800
Epoch 17/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
079 - acc: 0.9972 - val_loss: 0.0969 - val_acc: 0.9804
Epoch 18/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
087 - acc: 0.9970 - val_loss: 0.0962 - val_acc: 0.9804
Epoch 19/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
050 - acc: 0.9982 - val_loss: 0.0911 - val_acc: 0.9809
Epoch 20/20
60000/60000 [=====] - 2s 31us/step - loss: 0.0
086 - acc: 0.9970 - val_loss: 0.1027 - val_acc: 0.9801

```

```

In [63]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy


```

```
# for each key in history.history we will have a list of length equal  
to number of epochs
```

```
vy = history.history['val_loss']  
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10270734503133722

Test accuracy: 0.9801




```
In [64]: w_after = model_relu.get_weights()  
  
h1_w = w_after[0].flatten().reshape(-1,1)  
h2_w = w_after[2].flatten().reshape(-1,1)  
out_w = w_after[4].flatten().reshape(-1,1)  
  
fig = plt.figure()  
plt.title("Weight matrices after model trained")  
plt.subplot(1, 3, 1)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h1_w,color='b')
```

```
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



(b)MLP + Batch-Norm on hidden Layers + AdamOptimizer +RELU

```
In [65]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy thi
```

```

s condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.041 \Rightarrow N(0,\sigma) = N(0,0.041)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.058 \Rightarrow N(0,\sigma) = N(0,0.058)$ 
# out =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.098 \Rightarrow N(0,\sigma) = N(0,0.98)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(396, activation='relu', input_shape=(input_dim,),
    kernel_initializer=RandomNormal(mean=0.0, stddev=0.041, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(198, activation='relu', kernel_initializer=Random
Normal(mean=0.0, stddev=0.58, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
dense_47 (Dense)	(None, 396)	310860
batch_normalization_9 (Batch Normalization)	(None, 396)	1584
dense_48 (Dense)	(None, 198)	78606
batch_normalization_10 (Batch Normalization)	(None, 198)	792
dense_49 (Dense)	(None, 10)	1990
Total params: 393,832		
Trainable params: 392,644		
Non-trainable params: 1,188		


```
In [66]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 5s 78us/step - loss: 0.1936 - acc: 0.9422 - val_loss: 0.1086 - val_acc: 0.9683
Epoch 2/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0754 - acc: 0.9774 - val_loss: 0.0879 - val_acc: 0.9732
Epoch 3/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0486 - acc: 0.9851 - val_loss: 0.0773 - val_acc: 0.9780
Epoch 4/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0344 - acc: 0.9895 - val_loss: 0.0825 - val_acc: 0.9768
Epoch 5/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0253 - acc: 0.9922 - val_loss: 0.0682 - val_acc: 0.9795
Epoch 6/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0198 - acc: 0.9940 - val_loss: 0.0832 - val_acc: 0.9756
Epoch 7/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0150 - acc: 0.9957 - val_loss: 0.0825 - val_acc: 0.9768
Epoch 8/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0151 - acc: 0.9951 - val_loss: 0.0977 - val_acc: 0.9721
Epoch 9/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0162 - acc: 0.9946 - val_loss: 0.0842 - val_acc: 0.9761
Epoch 10/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0119 - acc: 0.9963 - val_loss: 0.0856 - val_acc: 0.9785
Epoch 11/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0102 - acc: 0.9969 - val_loss: 0.0772 - val_acc: 0.9805
```

```

Epoch 12/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0
083 - acc: 0.9974 - val_loss: 0.0779 - val_acc: 0.9793
Epoch 13/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0
106 - acc: 0.9964 - val_loss: 0.0855 - val_acc: 0.9767
Epoch 14/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0
105 - acc: 0.9965 - val_loss: 0.0938 - val_acc: 0.9762
Epoch 15/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0
100 - acc: 0.9967 - val_loss: 0.0899 - val_acc: 0.9781
Epoch 16/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0
061 - acc: 0.9981 - val_loss: 0.0896 - val_acc: 0.9792
Epoch 17/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0
069 - acc: 0.9976 - val_loss: 0.0796 - val_acc: 0.9812
Epoch 18/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0
072 - acc: 0.9978 - val_loss: 0.0913 - val_acc: 0.9780
Epoch 19/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0
069 - acc: 0.9976 - val_loss: 0.0923 - val_acc: 0.9803
Epoch 20/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0
072 - acc: 0.9977 - val_loss: 0.0831 - val_acc: 0.9795

```

```

In [67]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())

```

```
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08313359864633312

Test accuracy: 0.9795



In [68]: w_after = model_batch.get_weights()

```


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



(c) MLP + Dropout + AdamOptimizer +Relu+BatchNormalization

```
In [69]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batch-normalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(396, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.041, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.58, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_50 (Dense)	(None, 396)	310860
batch_normalization_11 (Batch Normalization)	(None, 396)	1584
dropout_11 (Dropout)	(None, 396)	0
dense_51 (Dense)	(None, 198)	78606
batch_normalization_12 (Batch Normalization)	(None, 198)	792
dropout_12 (Dropout)	(None, 198)	0

dense_52 (Dense)	(None, 10)	1990
------------------	------------	------

```

Total params: 393,832
Trainable params: 392,644
Non-trainable params: 1,188

```

```
In [70]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 83us/step - loss: 0.5034 - acc: 0.8470 - val_loss: 0.1727 - val_acc: 0.9459

Epoch 2/20

60000/60000 [=====] - 3s 50us/step - loss: 0.2593 - acc: 0.9214 - val_loss: 0.1276 - val_acc: 0.9596

Epoch 3/20

60000/60000 [=====] - 3s 51us/step - loss: 0.2024 - acc: 0.9383 - val_loss: 0.1144 - val_acc: 0.9643

Epoch 4/20

60000/60000 [=====] - 3s 50us/step - loss: 0.1763 - acc: 0.9468 - val_loss: 0.1017 - val_acc: 0.9673

Epoch 5/20

60000/60000 [=====] - 3s 50us/step - loss: 0.1585 - acc: 0.9516 - val_loss: 0.0922 - val_acc: 0.9711

Epoch 6/20

60000/60000 [=====] - 3s 51us/step - loss: 0.1445 - acc: 0.9550 - val_loss: 0.0894 - val_acc: 0.9727

Epoch 7/20

60000/60000 [=====] - 3s 50us/step - loss: 0.1318 - acc: 0.9594 - val_loss: 0.0817 - val_acc: 0.9752

Epoch 8/20

60000/60000 [=====] - 3s 50us/step - loss: 0.1251 - acc: 0.9614 - val_loss: 0.0756 - val_acc: 0.9761

```

Epoch 9/20
60000/60000 [=====] - 3s 50us/step - loss: 0.1
146 - acc: 0.9651 - val_loss: 0.0754 - val_acc: 0.9774
Epoch 10/20
60000/60000 [=====] - 3s 51us/step - loss: 0.1
115 - acc: 0.9658 - val_loss: 0.0694 - val_acc: 0.9788
Epoch 11/20
60000/60000 [=====] - 3s 50us/step - loss: 0.1
044 - acc: 0.9675 - val_loss: 0.0690 - val_acc: 0.9788
Epoch 12/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0
961 - acc: 0.9699 - val_loss: 0.0685 - val_acc: 0.9793
Epoch 13/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0
929 - acc: 0.9713 - val_loss: 0.0643 - val_acc: 0.9790
Epoch 14/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0
924 - acc: 0.9715 - val_loss: 0.0648 - val_acc: 0.9796
Epoch 15/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0
852 - acc: 0.9737 - val_loss: 0.0653 - val_acc: 0.9798
Epoch 16/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0
831 - acc: 0.9737 - val_loss: 0.0694 - val_acc: 0.9782
Epoch 17/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0
778 - acc: 0.9756 - val_loss: 0.0643 - val_acc: 0.9805
Epoch 18/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0
756 - acc: 0.9761 - val_loss: 0.0662 - val_acc: 0.9809
Epoch 19/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0
732 - acc: 0.9764 - val_loss: 0.0657 - val_acc: 0.9815
Epoch 20/20
60000/60000 [=====] - 3s 53us/step - loss: 0.0
739 - acc: 0.9766 - val_loss: 0.0672 - val_acc: 0.9810

```

```

In [71]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])

```

```

print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06723630262503284
Test accuracy: 0.981



```
In [72]: w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ' )
plt.show()
```



(d)MLP + Dropout + AdamOptimizer +Relu

```
In [73]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batch-normalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(396, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.071, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=None)) )
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))
```

```
model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_53 (Dense)	(None, 396)	310860
dropout_13 (Dropout)	(None, 396)	0
dense_54 (Dense)	(None, 198)	78606
dropout_14 (Dropout)	(None, 198)	0
dense_55 (Dense)	(None, 10)	1990

=====
Total params: 391,456
Trainable params: 391,456
Non-trainable params: 0
=====

```
In [74]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 59us/step - loss: 0.5

392 - acc: 0.8326 - val_loss: 0.1643 - val_acc: 0.9499

Epoch 2/20

60000/60000 [=====] - 2s 34us/step - loss: 0.2

342 - acc: 0.9297 - val_loss: 0.1232 - val_acc: 0.9626

Epoch 3/20

60000/60000 [=====] - 2s 34us/step - loss: 0.1

814 - acc: 0.9465 - val_loss: 0.1045 - val_acc: 0.9662

Epoch 4/20

60000/60000 [=====] - 2s 34us/step - loss: 0.1

```
540 - acc: 0.9553 - val_loss: 0.0877 - val_acc: 0.9716
Epoch 5/20
60000/60000 [=====] - 2s 34us/step - loss: 0.1
323 - acc: 0.9598 - val_loss: 0.0812 - val_acc: 0.9745
Epoch 6/20
60000/60000 [=====] - 2s 34us/step - loss: 0.1
158 - acc: 0.9652 - val_loss: 0.0742 - val_acc: 0.9749
Epoch 7/20
60000/60000 [=====] - 2s 34us/step - loss: 0.1
055 - acc: 0.9674 - val_loss: 0.0783 - val_acc: 0.9762
Epoch 8/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
979 - acc: 0.9706 - val_loss: 0.0711 - val_acc: 0.9789
Epoch 9/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
938 - acc: 0.9718 - val_loss: 0.0678 - val_acc: 0.9792
Epoch 10/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
863 - acc: 0.9734 - val_loss: 0.0658 - val_acc: 0.9798
Epoch 11/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
799 - acc: 0.9755 - val_loss: 0.0671 - val_acc: 0.9811
Epoch 12/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0
768 - acc: 0.9760 - val_loss: 0.0677 - val_acc: 0.9818
Epoch 13/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0
737 - acc: 0.9774 - val_loss: 0.0658 - val_acc: 0.9814
Epoch 14/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0
711 - acc: 0.9783 - val_loss: 0.0667 - val_acc: 0.9813
Epoch 15/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0
647 - acc: 0.9798 - val_loss: 0.0627 - val_acc: 0.9833
Epoch 16/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0
650 - acc: 0.9802 - val_loss: 0.0656 - val_acc: 0.9818
Epoch 17/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
```

```

615 - acc: 0.9807 - val_loss: 0.0637 - val_acc: 0.9827
Epoch 18/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
613 - acc: 0.9811 - val_loss: 0.0621 - val_acc: 0.9826
Epoch 19/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
587 - acc: 0.9817 - val_loss: 0.0612 - val_acc: 0.9824
Epoch 20/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0
557 - acc: 0.9820 - val_loss: 0.0617 - val_acc: 0.9822

```

```

In [75]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06173285202749103

Test accuracy: 0.9822



```
In [76]: w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2.Using hidden three layers

(a) MLP + ReLU + ADAM

```
In [99]: model_relu = Sequential()
model_relu.add(Dense(392, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.071, seed=None)))
model_relu.add(Dense(196, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=None)) )
model_relu.add(Dense(98, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.14, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
```

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_88 (Dense)	(None, 392)	307720
dense_89 (Dense)	(None, 196)	77028
dense_90 (Dense)	(None, 98)	19306
dense_91 (Dense)	(None, 10)	990

Total params: 405,044

Trainable params: 405,044

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 78us/step - loss: 0.2327 - acc: 0.9298 - val_loss: 0.1094 - val_acc: 0.9663

Epoch 2/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0863 - acc: 0.9734 - val_loss: 0.0958 - val_acc: 0.9712

Epoch 3/20

60000/60000 [=====] - 2s 36us/step - loss: 0.0572 - acc: 0.9818 - val_loss: 0.0762 - val_acc: 0.9771

Epoch 4/20

60000/60000 [=====] - 2s 36us/step - loss: 0.0389 - acc: 0.9874 - val_loss: 0.0765 - val_acc: 0.9775

Epoch 5/20

60000/60000 [=====] - 2s 36us/step - loss: 0.0280 - acc: 0.9909 - val_loss: 0.0714 - val_acc: 0.9789

Epoch 6/20


```
60000/60000 [=====] - 2s 36us/step - loss: 0.0
250 - acc: 0.9918 - val_loss: 0.0748 - val_acc: 0.9788
Epoch 7/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
231 - acc: 0.9922 - val_loss: 0.0820 - val_acc: 0.9775
Epoch 8/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
183 - acc: 0.9941 - val_loss: 0.0848 - val_acc: 0.9784
Epoch 9/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
177 - acc: 0.9942 - val_loss: 0.0875 - val_acc: 0.9801
Epoch 10/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
171 - acc: 0.9946 - val_loss: 0.0866 - val_acc: 0.9775
Epoch 11/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
139 - acc: 0.9952 - val_loss: 0.0837 - val_acc: 0.9803
Epoch 12/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
145 - acc: 0.9956 - val_loss: 0.0880 - val_acc: 0.9797
Epoch 13/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
122 - acc: 0.9959 - val_loss: 0.1167 - val_acc: 0.9753
Epoch 14/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
094 - acc: 0.9969 - val_loss: 0.0995 - val_acc: 0.9785
Epoch 15/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
116 - acc: 0.9963 - val_loss: 0.1228 - val_acc: 0.9756
Epoch 16/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
135 - acc: 0.9958 - val_loss: 0.0912 - val_acc: 0.9795
Epoch 17/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
101 - acc: 0.9967 - val_loss: 0.0883 - val_acc: 0.9801
Epoch 18/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0
093 - acc: 0.9970 - val_loss: 0.1071 - val_acc: 0.9766
Epoch 19/20
```

```
60000/60000 [=====] - 2s 36us/step - loss: 0.0096 - acc: 0.9968 - val_loss: 0.0926 - val_acc: 0.9809
Epoch 20/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0084 - acc: 0.9972 - val_loss: 0.1020 - val_acc: 0.9810
```

```
In [100]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
Test score: 0.10204988605169434
Test accuracy: 0.981
```



```
In [107]: w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
```

```
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



(b) MLP + Batch-Norm on hidden Layers + AdamOptimizer +RELU

```
In [144]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.041 \Rightarrow N(0, \sigma) = N(0, 0.041)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.058 \Rightarrow N(0, \sigma) = N(0, 0.058)$ 
# out =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.098 \Rightarrow N(0, \sigma) = N(0, 0.098)$ 

from keras.layers.normalization import BatchNormalization
```

```

model_batch = Sequential()

model_batch.add(Dense(396, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.041, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(198, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.58, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(98, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.08, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_124 (Dense)	(None, 396)	310860
batch_normalization_34 (Batch Normalization)	(None, 396)	1584
dense_125 (Dense)	(None, 198)	78606
batch_normalization_35 (Batch Normalization)	(None, 198)	792
dense_126 (Dense)	(None, 98)	19502
batch_normalization_36 (Batch Normalization)	(None, 98)	392
dense_127 (Dense)	(None, 10)	990
=====	=====	=====
Total params: 412,726		
Trainable params: 411,342		

Non-trainable params: 1,384

```
In [145]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy',  
metrics=['accuracy'])  
  
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 134us/step - loss: 0.2091 - acc: 0.9375 - val_loss: 0.1060 - val_acc: 0.9699

Epoch 2/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0775 - acc: 0.9766 - val_loss: 0.0932 - val_acc: 0.9693

Epoch 3/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0520 - acc: 0.9835 - val_loss: 0.0760 - val_acc: 0.9758

Epoch 4/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0370 - acc: 0.9884 - val_loss: 0.0805 - val_acc: 0.9743

Epoch 5/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0315 - acc: 0.9896 - val_loss: 0.0844 - val_acc: 0.9744

Epoch 6/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0259 - acc: 0.9915 - val_loss: 0.0762 - val_acc: 0.9775

Epoch 7/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0220 - acc: 0.9926 - val_loss: 0.0658 - val_acc: 0.9796

Epoch 8/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0174 - acc: 0.9941 - val_loss: 0.0707 - val_acc: 0.9784

Epoch 9/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0155 - acc: 0.9951 - val_loss: 0.0702 - val_acc: 0.9791

Epoch 10/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0

```

162 - acc: 0.9948 - val_loss: 0.0767 - val_acc: 0.9786
Epoch 11/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
139 - acc: 0.9952 - val_loss: 0.0880 - val_acc: 0.9756
Epoch 12/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
125 - acc: 0.9956 - val_loss: 0.0745 - val_acc: 0.9803
Epoch 13/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
107 - acc: 0.9967 - val_loss: 0.0843 - val_acc: 0.9785
Epoch 14/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
102 - acc: 0.9968 - val_loss: 0.0814 - val_acc: 0.9782
Epoch 15/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
129 - acc: 0.9956 - val_loss: 0.0857 - val_acc: 0.9772
Epoch 16/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
123 - acc: 0.9960 - val_loss: 0.0942 - val_acc: 0.9764
Epoch 17/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
085 - acc: 0.9972 - val_loss: 0.0931 - val_acc: 0.9768
Epoch 18/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
097 - acc: 0.9967 - val_loss: 0.0775 - val_acc: 0.9814
Epoch 19/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
075 - acc: 0.9977 - val_loss: 0.0797 - val_acc: 0.9811
Epoch 20/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0
083 - acc: 0.9972 - val_loss: 0.0866 - val_acc: 0.9790

```

```

In [146]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
          print('Test score:', score[0])
          print('Test accuracy:', score[1])

          fig,ax = plt.subplots(1,1)
          ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08658756976571458

Test accuracy: 0.979



```
In [147]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
```

```

ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



(c)MLP + Dropout + AdamOptimizer +Relu +batch

```

In [148]: model_drop = Sequential()

model_drop.add(Dense(396, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.041, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.58, seed=None)) )
model_drop.add(BatchNormalization())

```

```

model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.08, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_128 (Dense)	(None, 396)	310860
batch_normalization_37 (Batch Normalization)	(None, 396)	1584
dropout_29 (Dropout)	(None, 396)	0
dense_129 (Dense)	(None, 198)	78606
batch_normalization_38 (Batch Normalization)	(None, 198)	792
dropout_30 (Dropout)	(None, 198)	0
dense_130 (Dense)	(None, 198)	39402
batch_normalization_39 (Batch Normalization)	(None, 198)	792
dropout_31 (Dropout)	(None, 198)	0
dense_131 (Dense)	(None, 10)	1990
=====	=====	=====
Total params: 434,026		
Trainable params: 432,442		
Non-trainable params: 1,584		

```
In [149]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 154us/step - loss: 0.6811 - acc: 0.7915 - val_loss: 0.1994 - val_acc: 0.9372

Epoch 2/20

60000/60000 [=====] - 5s 77us/step - loss: 0.3135 - acc: 0.9063 - val_loss: 0.1488 - val_acc: 0.9541

Epoch 3/20

60000/60000 [=====] - 4s 68us/step - loss: 0.2522 - acc: 0.9262 - val_loss: 0.1256 - val_acc: 0.9613

Epoch 4/20

60000/60000 [=====] - 4s 69us/step - loss: 0.2187 - acc: 0.9357 - val_loss: 0.1158 - val_acc: 0.9633

Epoch 5/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1870 - acc: 0.9444 - val_loss: 0.1006 - val_acc: 0.9702

Epoch 6/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1704 - acc: 0.9490 - val_loss: 0.0914 - val_acc: 0.9721

Epoch 7/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1600 - acc: 0.9520 - val_loss: 0.0905 - val_acc: 0.9715

Epoch 8/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1457 - acc: 0.9568 - val_loss: 0.0898 - val_acc: 0.9740

Epoch 9/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1405 - acc: 0.9584 - val_loss: 0.0845 - val_acc: 0.9748

Epoch 10/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1341 - acc: 0.9605 - val_loss: 0.0814 - val_acc: 0.9757

Epoch 11/20

60000/60000 [=====] - 4s 68us/step - loss: 0.1254 - acc: 0.9626 - val_loss: 0.0740 - val_acc: 0.9782

```

Epoch 12/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1
228 - acc: 0.9636 - val_loss: 0.0763 - val_acc: 0.9777
Epoch 13/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1
138 - acc: 0.9651 - val_loss: 0.0729 - val_acc: 0.9785
Epoch 14/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1
068 - acc: 0.9680 - val_loss: 0.0701 - val_acc: 0.9804
Epoch 15/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1
060 - acc: 0.9682 - val_loss: 0.0664 - val_acc: 0.9812
Epoch 16/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1
009 - acc: 0.9700 - val_loss: 0.0733 - val_acc: 0.9796
Epoch 17/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0
940 - acc: 0.9703 - val_loss: 0.0721 - val_acc: 0.9798
Epoch 18/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0
925 - acc: 0.9719 - val_loss: 0.0657 - val_acc: 0.9808
Epoch 19/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0
875 - acc: 0.9735 - val_loss: 0.0643 - val_acc: 0.9813
Epoch 20/20
60000/60000 [=====] - 5s 76us/step - loss: 0.0
874 - acc: 0.9731 - val_loss: 0.0720 - val_acc: 0.9800

```

```

In [150]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())

```

```
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07200911457028705

Test accuracy: 0.98



In [151]: w_after = model_drop.get_weights()


```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



(d) MLP + Dropout + AdamOptimizer + Relu

```
In [114]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batch-normalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(396, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.071, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=None) ))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(98, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.14, seed=None) ))
model_drop.add(Dropout(0.5))
```



```
model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_96 (Dense)	(None, 396)	310860
dropout_21 (Dropout)	(None, 396)	0
dense_97 (Dense)	(None, 198)	78606
dropout_22 (Dropout)	(None, 198)	0
dense_98 (Dense)	(None, 98)	19502
dropout_23 (Dropout)	(None, 98)	0
dense_99 (Dense)	(None, 10)	990
Total params: 409,958		
Trainable params: 409,958		
Non-trainable params: 0		

```
In [115]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 82us/step - loss: 0.9453 - acc: 0.6996 - val_loss: 0.2284 - val_acc: 0.9359

Epoch 2/20

60000/60000 [=====] - 2s 39us/step - loss: 0.3

```
629 - acc: 0.8999 - val_loss: 0.1711 - val_acc: 0.9529
Epoch 3/20
60000/60000 [=====] - 2s 39us/step - loss: 0.2
708 - acc: 0.9277 - val_loss: 0.1486 - val_acc: 0.9585
Epoch 4/20
60000/60000 [=====] - 2s 39us/step - loss: 0.2
256 - acc: 0.9383 - val_loss: 0.1244 - val_acc: 0.9637
Epoch 5/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
934 - acc: 0.9485 - val_loss: 0.1097 - val_acc: 0.9702
Epoch 6/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
753 - acc: 0.9518 - val_loss: 0.1028 - val_acc: 0.9732
Epoch 7/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
570 - acc: 0.9579 - val_loss: 0.0989 - val_acc: 0.9739
Epoch 8/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
467 - acc: 0.9608 - val_loss: 0.0924 - val_acc: 0.9750
Epoch 9/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
378 - acc: 0.9621 - val_loss: 0.0874 - val_acc: 0.9752
Epoch 10/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
238 - acc: 0.9666 - val_loss: 0.0862 - val_acc: 0.9773
Epoch 11/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
163 - acc: 0.9675 - val_loss: 0.0867 - val_acc: 0.9755
Epoch 12/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
129 - acc: 0.9689 - val_loss: 0.0824 - val_acc: 0.9792
Epoch 13/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
062 - acc: 0.9712 - val_loss: 0.0858 - val_acc: 0.9789
Epoch 14/20
60000/60000 [=====] - 2s 39us/step - loss: 0.1
023 - acc: 0.9722 - val_loss: 0.0827 - val_acc: 0.9775
Epoch 15/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0
```

```

988 - acc: 0.9732 - val_loss: 0.0854 - val_acc: 0.9784
Epoch 16/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0
936 - acc: 0.9735 - val_loss: 0.0804 - val_acc: 0.9808
Epoch 17/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0
918 - acc: 0.9744 - val_loss: 0.0789 - val_acc: 0.9785
Epoch 18/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0
863 - acc: 0.9751 - val_loss: 0.0776 - val_acc: 0.9794
Epoch 19/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0
880 - acc: 0.9755 - val_loss: 0.0754 - val_acc: 0.9805
Epoch 20/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0
827 - acc: 0.9774 - val_loss: 0.0780 - val_acc: 0.9806

```

```

In [116]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy


```

```
# for each key in history.history we will have a list of length equal  
to number of epochs
```

```
vy = history.history['val_loss']  
ty = history.history['loss']  
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0779710623358118

Test accuracy: 0.9806




```
In [118]: w_after = model_drop.get_weights()  
  
h1_w = w_after[0].flatten().reshape(-1,1)  
h2_w = w_after[2].flatten().reshape(-1,1)  
h3_w = w_after[2].flatten().reshape(-1,1)  
out_w = w_after[4].flatten().reshape(-1,1)  
  
fig = plt.figure()  
plt.title("Weight matrices after model trained")  
plt.subplot(1, 4, 1)  
plt.title("Trained model Weights")  
ax = sns.violinplot(y=h1_w,color='b')
```

```
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3.using five layers

(a)MLP + ReLU + ADAM

```
In [119]: model_relu = Sequential()
model_relu.add(Dense(392, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.071, seed=None)))
model_relu.add(Dense(196, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=None)) )
model_relu.add(Dense(98, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.14, seed=None)) )
model_relu.add(Dense(49, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.2, seed=None)) )
model_relu.add(Dense(25, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.28, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_100 (Dense)	(None, 392)	307720
dense_101 (Dense)	(None, 196)	77028
dense_102 (Dense)	(None, 98)	19306
dense_103 (Dense)	(None, 49)	4851
dense_104 (Dense)	(None, 25)	1250
dense_105 (Dense)	(None, 10)	260
=====	=====	=====
Total params: 410,415		

Trainable params: 410,415

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 85us/step - loss: 0.3

188 - acc: 0.9015 - val_loss: 0.1328 - val_acc: 0.9596

Epoch 2/20

60000/60000 [=====] - 2s 41us/step - loss: 0.1

040 - acc: 0.9690 - val_loss: 0.1062 - val_acc: 0.9692

Epoch 3/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0

664 - acc: 0.9793 - val_loss: 0.0807 - val_acc: 0.9741

Epoch 4/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0

523 - acc: 0.9839 - val_loss: 0.0929 - val_acc: 0.9722

Epoch 5/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0

410 - acc: 0.9868 - val_loss: 0.0820 - val_acc: 0.9759

Epoch 6/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0

312 - acc: 0.9901 - val_loss: 0.0800 - val_acc: 0.9764

Epoch 7/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0

303 - acc: 0.9903 - val_loss: 0.0887 - val_acc: 0.9743

Epoch 8/20

60000/60000 [=====] - 2s 40us/step - loss: 0.0

273 - acc: 0.9907 - val_loss: 0.1054 - val_acc: 0.9743

Epoch 9/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0

218 - acc: 0.9929 - val_loss: 0.0935 - val_acc: 0.9781

Epoch 10/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0

259 - acc: 0.9915 - val_loss: 0.0956 - val_acc: 0.9767

Epoch 11/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0

208 - acc: 0.9932 - val_loss: 0.0886 - val_acc: 0.9782

Epoch 12/20

```

60000/60000 [=====] - 2s 40us/step - loss: 0.0
147 - acc: 0.9954 - val_loss: 0.0851 - val_acc: 0.9807
Epoch 13/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0
150 - acc: 0.9953 - val_loss: 0.0986 - val_acc: 0.9753
Epoch 14/20
60000/60000 [=====] - 2s 40us/step - loss: 0.0
172 - acc: 0.9944 - val_loss: 0.0848 - val_acc: 0.9810
Epoch 15/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0
167 - acc: 0.9950 - val_loss: 0.0950 - val_acc: 0.9782
Epoch 16/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0
132 - acc: 0.9956 - val_loss: 0.0940 - val_acc: 0.9808
Epoch 17/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0
169 - acc: 0.9945 - val_loss: 0.0892 - val_acc: 0.9796
Epoch 18/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0
087 - acc: 0.9972 - val_loss: 0.0826 - val_acc: 0.9799
Epoch 19/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0
140 - acc: 0.9958 - val_loss: 0.0925 - val_acc: 0.9818
Epoch 20/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0
098 - acc: 0.9968 - val_loss: 0.1038 - val_acc: 0.9776

```

```

In [120]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])

```



```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10375401149833269

Test accuracy: 0.9776



```
In [127]: w_after = model_relu.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1,6 ,5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1,6,6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ' )
plt.show()
```



(b)MLP + Batch-Norm on hidden Layers + AdamOptimizer +RELU

```
In [136]: # Multilayer perceptron

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(392, activation='relu', input_shape=(input_dim,),
    kernel_initializer=RandomNormal(mean=0.0, stddev=0.041, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(196, activation='relu', kernel_initializer=Random
Normal(mean=0.0, stddev=0.58, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(98, activation='relu', kernel_initializer=RandomN
```

```

ormal(mean=0.0, stddev=0.08, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(49, activation='relu', kernel_initializer=RandomN
ormal(mean=0.0, stddev=0.11, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(25, activation='relu', kernel_initializer=RandomN
ormal(mean=0.0, stddev=0.16, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_112 (Dense)	(None, 392)	307720
batch_normalization_29 (Batch Normalization)	(None, 392)	1568
dense_113 (Dense)	(None, 196)	77028
batch_normalization_30 (Batch Normalization)	(None, 196)	784
dense_114 (Dense)	(None, 98)	19306
batch_normalization_31 (Batch Normalization)	(None, 98)	392
dense_115 (Dense)	(None, 49)	4851
batch_normalization_32 (Batch Normalization)	(None, 49)	196
dense_116 (Dense)	(None, 25)	1250
batch_normalization_33 (Batch Normalization)	(None, 25)	100

```
dense_117 (Dense)                (None, 10)                260
=====
Total params: 413,455
Trainable params: 411,935
Non-trainable params: 1,520
```

```
In [137]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 9s 152us/step - loss: 0.2931 - acc: 0.9211 - val_loss: 0.1364 - val_acc: 0.9580
Epoch 2/20
60000/60000 [=====] - 6s 93us/step - loss: 0.1012 - acc: 0.9707 - val_loss: 0.1037 - val_acc: 0.9693
Epoch 3/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0665 - acc: 0.9798 - val_loss: 0.0901 - val_acc: 0.9724
Epoch 4/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0539 - acc: 0.9829 - val_loss: 0.0897 - val_acc: 0.9719
Epoch 5/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0414 - acc: 0.9867 - val_loss: 0.0849 - val_acc: 0.9738
Epoch 6/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0372 - acc: 0.9885 - val_loss: 0.0930 - val_acc: 0.9743
Epoch 7/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0307 - acc: 0.9904 - val_loss: 0.0927 - val_acc: 0.9748
Epoch 8/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0281 - acc: 0.9910 - val_loss: 0.1120 - val_acc: 0.9705
Epoch 9/20
```

```

60000/60000 [=====] - 5s 89us/step - loss: 0.0
247 - acc: 0.9921 - val_loss: 0.0848 - val_acc: 0.9759
Epoch 10/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
239 - acc: 0.9921 - val_loss: 0.0779 - val_acc: 0.9785
Epoch 11/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
225 - acc: 0.9926 - val_loss: 0.0817 - val_acc: 0.9779
Epoch 12/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
210 - acc: 0.9933 - val_loss: 0.0932 - val_acc: 0.9746
Epoch 13/20
60000/60000 [=====] - 6s 101us/step - loss: 0.
0189 - acc: 0.9939 - val_loss: 0.0800 - val_acc: 0.9801
Epoch 14/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
146 - acc: 0.9952 - val_loss: 0.0771 - val_acc: 0.9811
Epoch 15/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
151 - acc: 0.9952 - val_loss: 0.0880 - val_acc: 0.9786
Epoch 16/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
164 - acc: 0.9945 - val_loss: 0.0855 - val_acc: 0.9774
Epoch 17/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0
149 - acc: 0.9950 - val_loss: 0.0943 - val_acc: 0.9775
Epoch 18/20
60000/60000 [=====] - 6s 101us/step - loss: 0.
0133 - acc: 0.9956 - val_loss: 0.0930 - val_acc: 0.9783
Epoch 19/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0
118 - acc: 0.9963 - val_loss: 0.0843 - val_acc: 0.9799
Epoch 20/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0
112 - acc: 0.9961 - val_loss: 0.0919 - val_acc: 0.9799

```

```

In [138]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
          print('Test score:', score[0])
          print('Test accuracy:', score[1])

```

```

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.0918853635030333

Test accuracy: 0.9799



```
In [139]: w_after = model_batch.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```



```

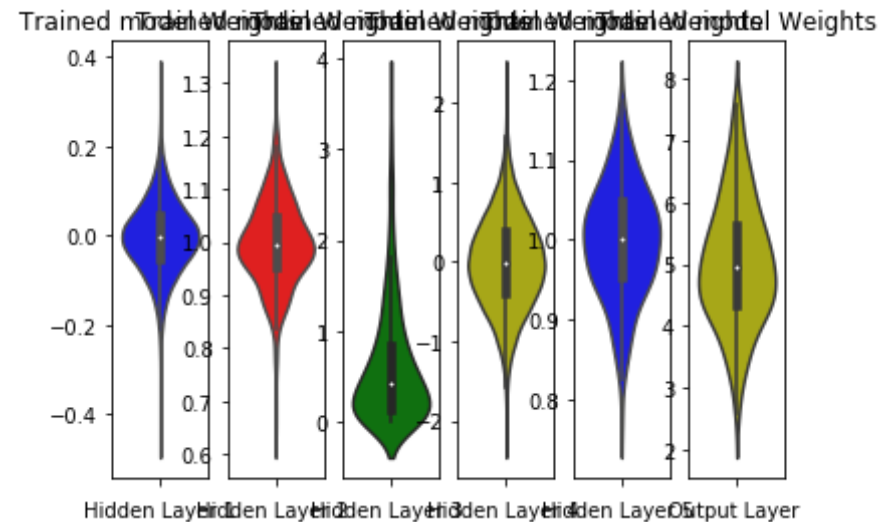
plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='y')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1,6 ,5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1,6,6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



(c)MLP + Dropout + AdamOptimizer +Relu +batch

```
In [152]: model_drop = Sequential()

model_drop.add(Dense(396, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.041, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomN
ormal(mean=0.0, stddev=0.58, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomN
ormal(mean=0.0, stddev=0.08, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomN
ormal(mean=0.0, stddev=0.11, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomN
ormal(mean=0.0, stddev=0.16, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_132 (Dense)	(None, 396)	310860

batch_normalization_40 (Batch Normalization)	(None, 396)	1584
dropout_32 (Dropout)	(None, 396)	0
dense_133 (Dense)	(None, 198)	78606
batch_normalization_41 (Batch Normalization)	(None, 198)	792
dropout_33 (Dropout)	(None, 198)	0
dense_134 (Dense)	(None, 198)	39402
batch_normalization_42 (Batch Normalization)	(None, 198)	792
dropout_34 (Dropout)	(None, 198)	0
dense_135 (Dense)	(None, 198)	39402
batch_normalization_43 (Batch Normalization)	(None, 198)	792
dropout_35 (Dropout)	(None, 198)	0
dense_136 (Dense)	(None, 198)	39402
batch_normalization_44 (Batch Normalization)	(None, 198)	792
dropout_36 (Dropout)	(None, 198)	0
dense_137 (Dense)	(None, 10)	1990

Total params: 514,414
 Trainable params: 512,038
 Non-trainable params: 2,376

```
In [153]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epoch
s=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 196us/step - loss: 1.2813 - acc: 0.5969 - val_loss: 0.3154 - val_acc: 0.9050

Epoch 2/20

60000/60000 [=====] - 6s 96us/step - loss: 0.4875 - acc: 0.8544 - val_loss: 0.1973 - val_acc: 0.9420

Epoch 3/20

60000/60000 [=====] - 6s 95us/step - loss: 0.3567 - acc: 0.8957 - val_loss: 0.1590 - val_acc: 0.9525

Epoch 4/20

60000/60000 [=====] - 6s 95us/step - loss: 0.2937 - acc: 0.9160 - val_loss: 0.1432 - val_acc: 0.9587

Epoch 5/20

60000/60000 [=====] - 6s 96us/step - loss: 0.2639 - acc: 0.9248 - val_loss: 0.1436 - val_acc: 0.9602

Epoch 6/20

60000/60000 [=====] - 6s 95us/step - loss: 0.2352 - acc: 0.9354 - val_loss: 0.1306 - val_acc: 0.9639

Epoch 7/20

60000/60000 [=====] - 6s 95us/step - loss: 0.2151 - acc: 0.9393 - val_loss: 0.1199 - val_acc: 0.9666

Epoch 8/20

60000/60000 [=====] - 6s 95us/step - loss: 0.2029 - acc: 0.9437 - val_loss: 0.1166 - val_acc: 0.9684

Epoch 9/20

60000/60000 [=====] - 6s 96us/step - loss: 0.1902 - acc: 0.9475 - val_loss: 0.1049 - val_acc: 0.9722

Epoch 10/20

60000/60000 [=====] - 6s 95us/step - loss: 0.1789 - acc: 0.9497 - val_loss: 0.1086 - val_acc: 0.9717

Epoch 11/20

60000/60000 [=====] - 6s 95us/step - loss: 0.1688 - acc: 0.9523 - val_loss: 0.1016 - val_acc: 0.9719

Epoch 12/20

60000/60000 [=====] - 6s 96us/step - loss: 0.1570 - acc: 0.9564 - val_loss: 0.0948 - val_acc: 0.9741

```

579 - acc: 0.9504 - val_loss: 0.0940 - val_acc: 0.9741
Epoch 13/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1
548 - acc: 0.9578 - val_loss: 0.0969 - val_acc: 0.9736
Epoch 14/20
60000/60000 [=====] - 6s 104us/step - loss: 0.
1479 - acc: 0.9587 - val_loss: 0.0883 - val_acc: 0.9763
Epoch 15/20
60000/60000 [=====] - 7s 109us/step - loss: 0.
1417 - acc: 0.9610 - val_loss: 0.0894 - val_acc: 0.9769
Epoch 16/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1
356 - acc: 0.9621 - val_loss: 0.0871 - val_acc: 0.9771
Epoch 17/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1
304 - acc: 0.9640 - val_loss: 0.0863 - val_acc: 0.9775
Epoch 18/20
60000/60000 [=====] - 6s 96us/step - loss: 0.1
303 - acc: 0.9640 - val_loss: 0.0852 - val_acc: 0.9769
Epoch 19/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1
226 - acc: 0.9661 - val_loss: 0.0843 - val_acc: 0.9775
Epoch 20/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1
194 - acc: 0.9671 - val_loss: 0.0834 - val_acc: 0.9772

```

```

In [154]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epo

```

```

chs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter vali
dation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal
to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.083420745628234
Test accuracy: 0.9772



```

In [155]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)

```

```

h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='y')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1,6 ,5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1,6,6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



(d) MLP + Dropout + AdamOptimizer +Relu

```
In [140]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batch-normalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(396, activation='relu', input_shape=(input_dim,),
kernel_initializer=RandomNormal(mean=0.0, stddev=0.071, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(198, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1, seed=None)) )
model_drop.add(Dropout(0.5))

model_drop.add(Dense(98, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.14, seed=None)) )
model_drop.add(Dropout(0.5))
```



```

model_drop.add(Dense(98, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.2, seed=None)) )
model_drop.add(Dropout(0.5))

model_drop.add(Dense(98, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.28, seed=None)) )
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```

Layer (type)	Output Shape	Param #
dense_118 (Dense)	(None, 396)	310860
dropout_24 (Dropout)	(None, 396)	0
dense_119 (Dense)	(None, 198)	78606
dropout_25 (Dropout)	(None, 198)	0
dense_120 (Dense)	(None, 98)	19502
dropout_26 (Dropout)	(None, 98)	0
dense_121 (Dense)	(None, 98)	9702
dropout_27 (Dropout)	(None, 98)	0
dense_122 (Dense)	(None, 98)	9702
dropout_28 (Dropout)	(None, 98)	0
dense_123 (Dense)	(None, 10)	990
Total params: 429,362		

Trainable params: 429,362

```
trainable params: 429,362  
Non-trainable params: 0
```

```
In [141]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', m  
etrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epoch  
s=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/20
```

```
60000/60000 [=====] - 7s 109us/step - loss: 3.  
0582 - acc: 0.1556 - val_loss: 1.9679 - val_acc: 0.2860
```

```
Epoch 2/20
```

```
60000/60000 [=====] - 3s 47us/step - loss: 1.7  
921 - acc: 0.3399 - val_loss: 1.4082 - val_acc: 0.4728
```

```
Epoch 3/20
```

```
60000/60000 [=====] - 3s 47us/step - loss: 1.4  
161 - acc: 0.4615 - val_loss: 1.1216 - val_acc: 0.5522
```

```
Epoch 4/20
```

```
60000/60000 [=====] - 3s 49us/step - loss: 1.2  
070 - acc: 0.5386 - val_loss: 0.9771 - val_acc: 0.6124
```

```
Epoch 5/20
```

```
60000/60000 [=====] - 3s 53us/step - loss: 1.0  
516 - acc: 0.6083 - val_loss: 0.7809 - val_acc: 0.7109
```

```
Epoch 6/20
```

```
60000/60000 [=====] - 3s 53us/step - loss: 0.9  
086 - acc: 0.6694 - val_loss: 0.6112 - val_acc: 0.7783
```

```
Epoch 7/20
```

```
60000/60000 [=====] - 3s 52us/step - loss: 0.7  
851 - acc: 0.7270 - val_loss: 0.5624 - val_acc: 0.7598
```

```
Epoch 8/20
```

```
60000/60000 [=====] - 3s 47us/step - loss: 0.6  
841 - acc: 0.7842 - val_loss: 0.5044 - val_acc: 0.8145
```

```
Epoch 9/20
```

```
60000/60000 [=====] - 3s 47us/step - loss: 0.5  
944 - acc: 0.8349 - val_loss: 0.3985 - val_acc: 0.8392
```

```
Epoch 10/20
```

```
60000/60000 [=====] - 3s 47us/step - loss: 0.5
```

```

267 - acc: 0.8578 - val_loss: 0.3241 - val_acc: 0.8919
Epoch 11/20
60000/60000 [=====] - 3s 47us/step - loss: 0.4
723 - acc: 0.8759 - val_loss: 0.3136 - val_acc: 0.8852
Epoch 12/20
60000/60000 [=====] - 3s 47us/step - loss: 0.4
247 - acc: 0.8891 - val_loss: 0.2491 - val_acc: 0.9310
Epoch 13/20
60000/60000 [=====] - 3s 47us/step - loss: 0.4
012 - acc: 0.8972 - val_loss: 0.2242 - val_acc: 0.9437
Epoch 14/20
60000/60000 [=====] - 3s 46us/step - loss: 0.3
728 - acc: 0.9068 - val_loss: 0.2165 - val_acc: 0.9465
Epoch 15/20
60000/60000 [=====] - 3s 47us/step - loss: 0.3
586 - acc: 0.9112 - val_loss: 0.2071 - val_acc: 0.9493
Epoch 16/20
60000/60000 [=====] - 3s 47us/step - loss: 0.3
419 - acc: 0.9165 - val_loss: 0.1928 - val_acc: 0.9538
Epoch 17/20
60000/60000 [=====] - 3s 46us/step - loss: 0.3
187 - acc: 0.9223 - val_loss: 0.1964 - val_acc: 0.9526
Epoch 18/20
60000/60000 [=====] - 3s 47us/step - loss: 0.3
109 - acc: 0.9246 - val_loss: 0.1897 - val_acc: 0.9560
Epoch 19/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2
971 - acc: 0.9295 - val_loss: 0.1947 - val_acc: 0.9548
Epoch 20/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2
931 - acc: 0.9297 - val_loss: 0.1779 - val_acc: 0.9567

```

```

In [142]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
          print('Test score:', score[0])
          print('Test accuracy:', score[1])

          fig,ax = plt.subplots(1,1)
          ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.17786330469548703

Test accuracy: 0.9567



```
In [143]: w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```

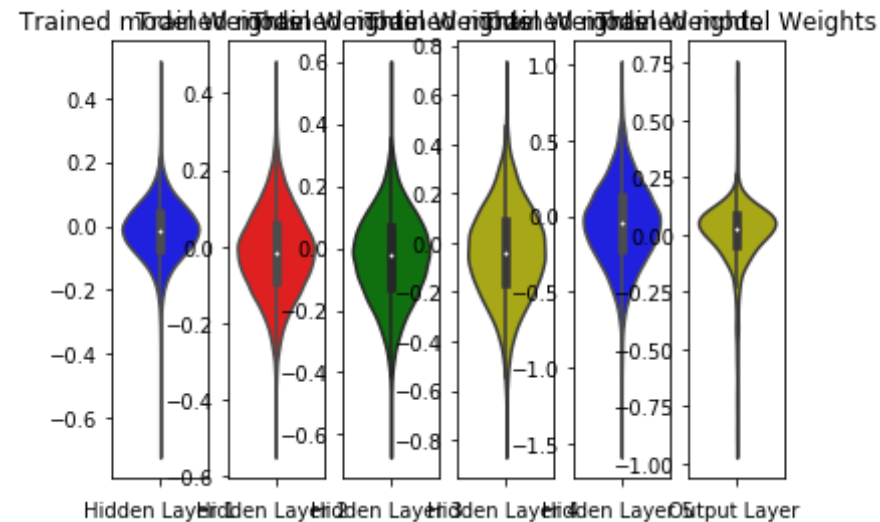
plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='y')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1,6 ,5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='b')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1,6,6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



conclusion

- STEP1:collected mnist data
- STEP2:converted into single dimension i.e column
- STEP3:using 2/3/5 hidden layers and activation function as RELU and optimization as ADAM and BATCHING and DROPOUT are done for the hidden layers

```
In [158]: from prettytable import PrettyTable

print("*****using 2 hidden layers*****")
names = ['mlp-relu-adam', 'mlp-batch-adam-relu', 'mlp-drop-adam-relu-batch', 'mlp-drop-adam-relu']

acc = [98.01,97.75,98.10,98.22]

numbering = [1,2,3,4]

# Initializing prettytable
table = PrettyTable()

# Adding columns
table.add_column("S.NO.", numbering)
table.add_column("MODEL", names)
table.add_column("acc ", acc)

# Printing the Table
print(table)

print("*****using 3 hidden layers*****")
names = ['mlp-relu-adam', 'mlp-batch-adam-relu', 'mlp-drop-adam-relu-batch', 'mlp-drop-adam-relu']

acc = [98.10,97.90,98.00,95.06]
```

```

numbering = [1,2,3,4]

# Initializing prettytable
table = PrettyTable()

# Adding columns
table.add_column("S.NO.", numbering)
table.add_column("MODEL", names)
table.add_column("acc ", acc)

# Printing the Table
print(table)

print("*****using 5 hidden layers*****")
names = ['mlp-relu-adam', 'mlp-batch-adam-relu', 'mlp-drop-adam-relu-batch', 'mlp-drop-adam-relu']

acc = [97.76, 97.99, 98.72, 95.67]

numbering = [1,2,3,4]

# Initializing prettytable
table = PrettyTable()

# Adding columns
table.add_column("S.NO.", numbering)
table.add_column("MODEL", names)
table.add_column("acc ", acc)

# Printing the Table
print(table)

*****using 2 hidden layers*****
*****
+-----+-----+-----+
| S.NO. | MODEL | acc |
+-----+-----+-----+
| 1 | mlp-relu-adam | 97.76 |
| 2 | mlp-batch-adam-relu | 97.99 |
| 3 | mlp-drop-adam-relu-batch | 98.72 |
| 4 | mlp-drop-adam-relu | 95.67 |
+-----+-----+-----+

```


1	mlp-relu-adam	98.01
2	mlp-batch-adam-relu	97.75
3	mlp-drop-adam-relu-batch	98.1
4	mlp-drop-adam-relu	98.22

-----+-----+-----+-----+
 *****using 3 hidden layers*****

S.NO.	MODEL	acc
1	mlp-relu-adam	98.1
2	mlp-batch-adam-relu	97.9
3	mlp-drop-adam-relu-batch	98.0
4	mlp-drop-adam-relu	95.06

-----+-----+-----+-----+
 *****using 5 hidden layers*****

S.NO.	MODEL	acc
1	mlp-relu-adam	97.76
2	mlp-batch-adam-relu	97.99
3	mlp-drop-adam-relu-batch	98.72
4	mlp-drop-adam-relu	95.67

-----+-----+-----+-----+